



INTRODUCING NOSQL AND MONGODB

Pom Bleeksma & Lucas Jellema
14 March 2017, Nieuwegein



mongoDB

AGENDA

INTRODUCTION

WHAT'S NOT SQL ABOUT
NOSQL – HISTORY,
BACKGROUND, USE CASES

OVERVIEW OF NOSQL
DATABASES

INTRODUCING MONGODB

GETTING STARTED WITH
MONGODB – INSTALLATION,
CONFIGURATION,
ADMINISTRATION




DEVELOPING WITH MONGODB
– JAVA & NODE.JS

HANDSON WORKSHOP



TRADITIONAL APPROACH: ALL DATA IN ENTERPRISE RELATIONAL DATABASE


Conventional view of
Data Management



TRADITIONAL APPROACH: WITHOUT MUCH REGARD FOR

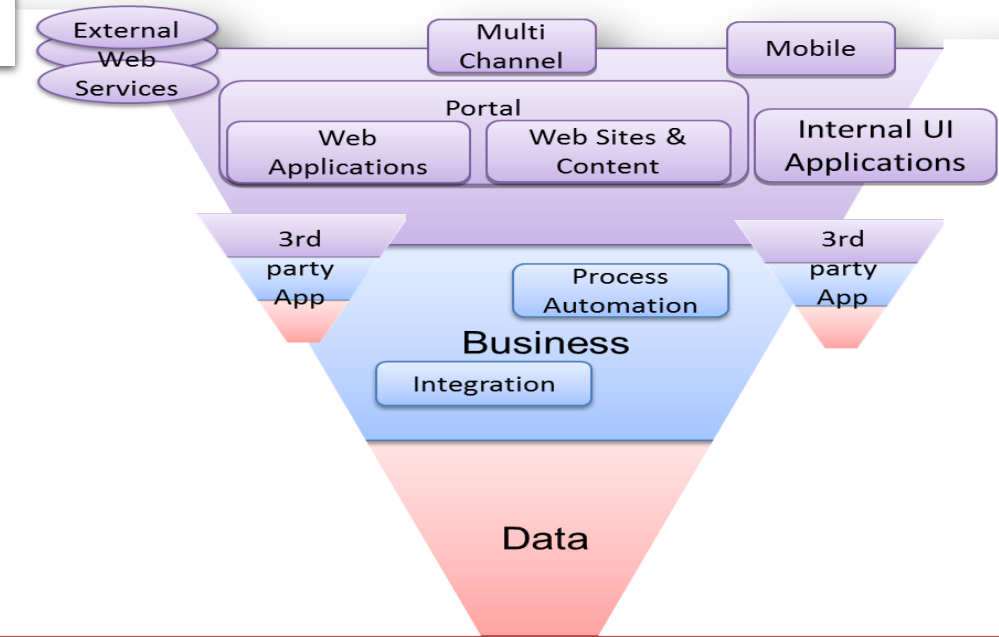
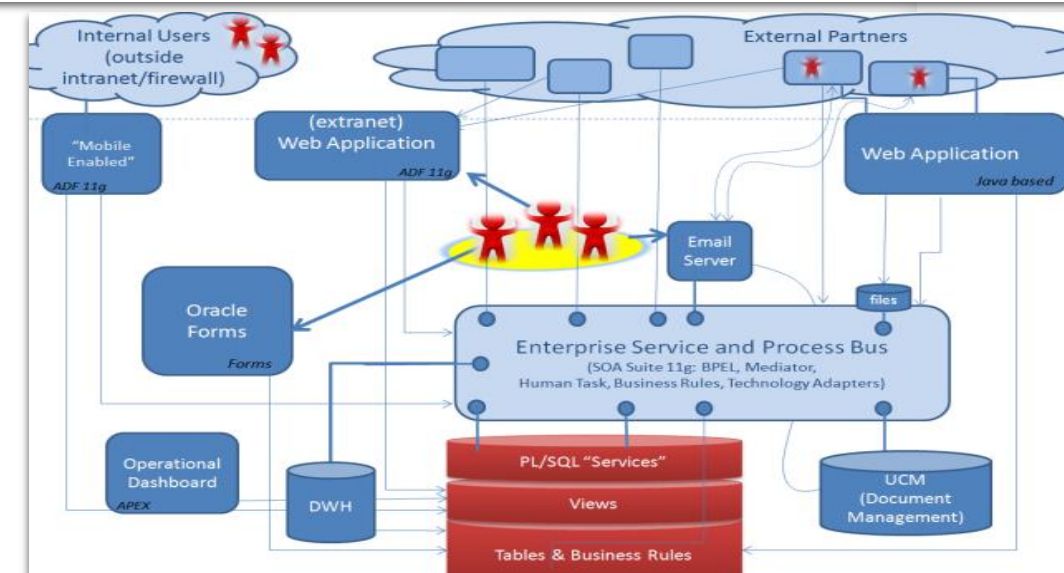

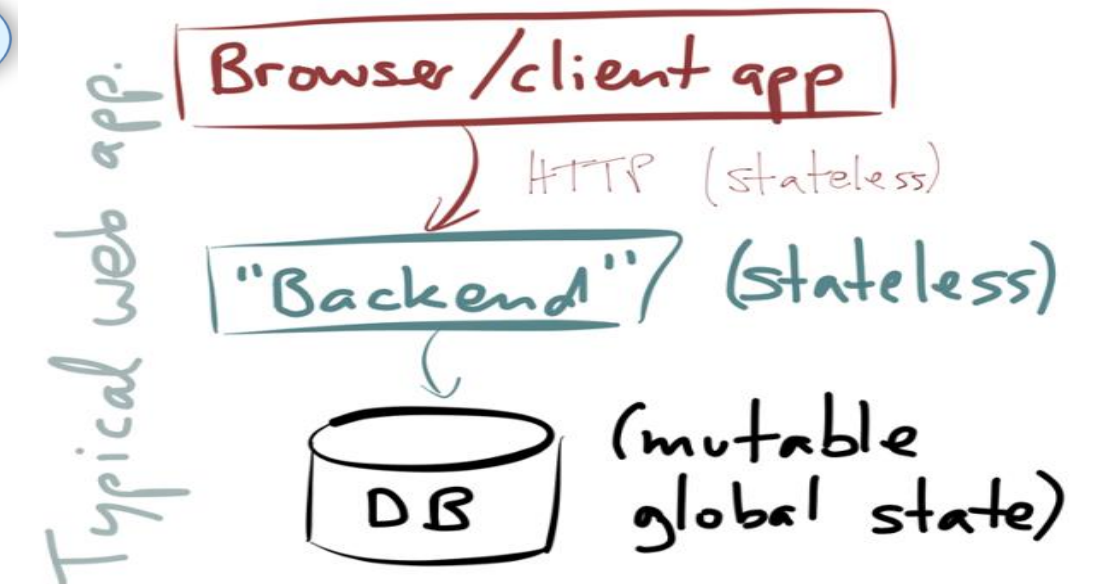

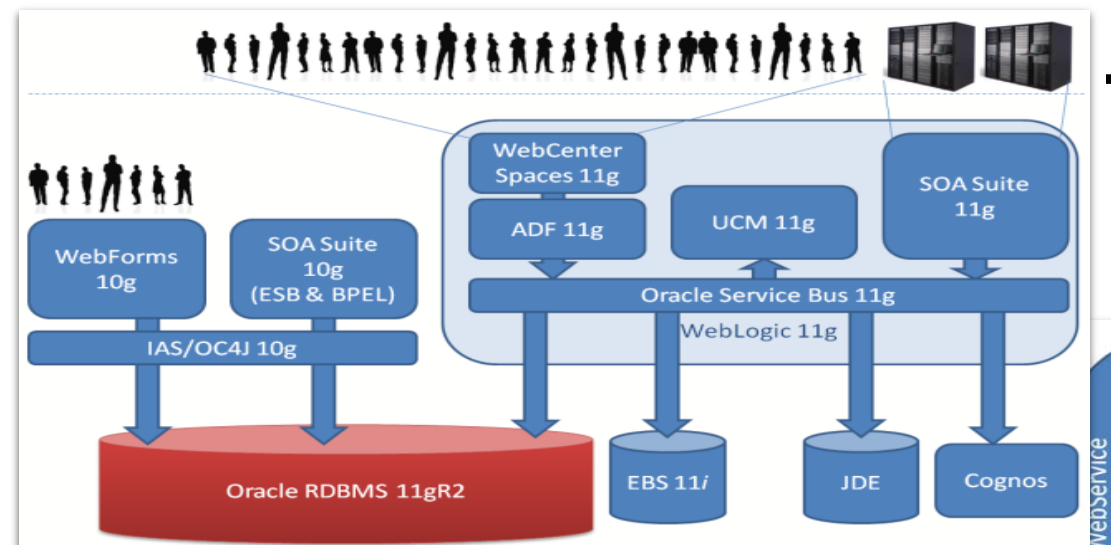
- What is it used for?
 - By whom, where, in what way, using which tools
- Ratio between writes and reads
- Ad hoc access?
- What format is it in/should it be in?
- What confidentiality & integrity is required?
- How much of it?
- How long relevant (hot vs cold vs dead)?
- How fine grained and how accurate?
- How much change?

Conventional view of
Data Management




DATABASE: SINGLE POINT OF ... SCALABILITY | AVAILABILITY | CONSISTENCY

- Single Point of ...



RELATIONAL DATABASES

- Based on **relational model** of data (E.F. Codd), a mathematical foundation
- Uses **SQL** for query, DML and DDL
- Transactions are **ACID** (Atomicity, Consistency, Isolation, Durability)
 - **All** or nothing
 - **Constraint** Compliant
 - **Individual** experience [in a multi-session environment] (aka concurrency)
 - **Down** does not hurt



THE HOLY GRAIL OF NORMALIZATION

- Normalize to prevent data redundancy, discrepancies (split brain) and storage waste
- Pragmatic approaches - recognizing the fact that some data is read far more frequently than that it is created and modified
 - Materialized View
 - Virtual Column
 - Index
 - Sharding
 - Replication
 - Data Guard
 - Data Warehouse & Data Marts
 - Client Side Caching
 - In Memory Database



ACID COMES AT A COST

- Transaction results have to be persisted in order to guarantee **D**
- Concurrency requires some degree of locking (and multi-versioning) in order to have **I**
- Constraint compliance (unique key, foreign key) means all data hangs together (as do all transactions) in order to have **C**



ACID COMES AT A COST

- ACID means:
 - All data closely together (and sometimes far from applications and users)
 - All transactions centrally organized
 - All data persisted as part of the transaction
- ACID means:
 - Hard to scale out (horizontally)
- Oracle offers scale up with Real Application Clusters



THE RELATIONAL MODEL IN PRACTICE


- Traditional Relational Data Model has severe impact on physical disk performance
 - Transaction Log => Sequential Write (append to file)
 - Data Blocks require much more expensive Random Access disk writes
- Indexes (B-Tree, Bitmap, ...) are used to speed up query (read) performance
 - and slow down transactions
- Relational data does not [always] map naturally to data format required in application (OO, JSON, XML)
- Capability to join and construct ad-hoc queries across the entire data model is powerful
- Declarative integrity constraints allow for strict enforcement of data quality rules
 - *“the data may be non sensical, but at least it adheres to the rules, whatever the origin”*
- SQL has been embraced by huge numbers of tools & technologies and IT/Data professionals

DATABASES RE-EVALUATED

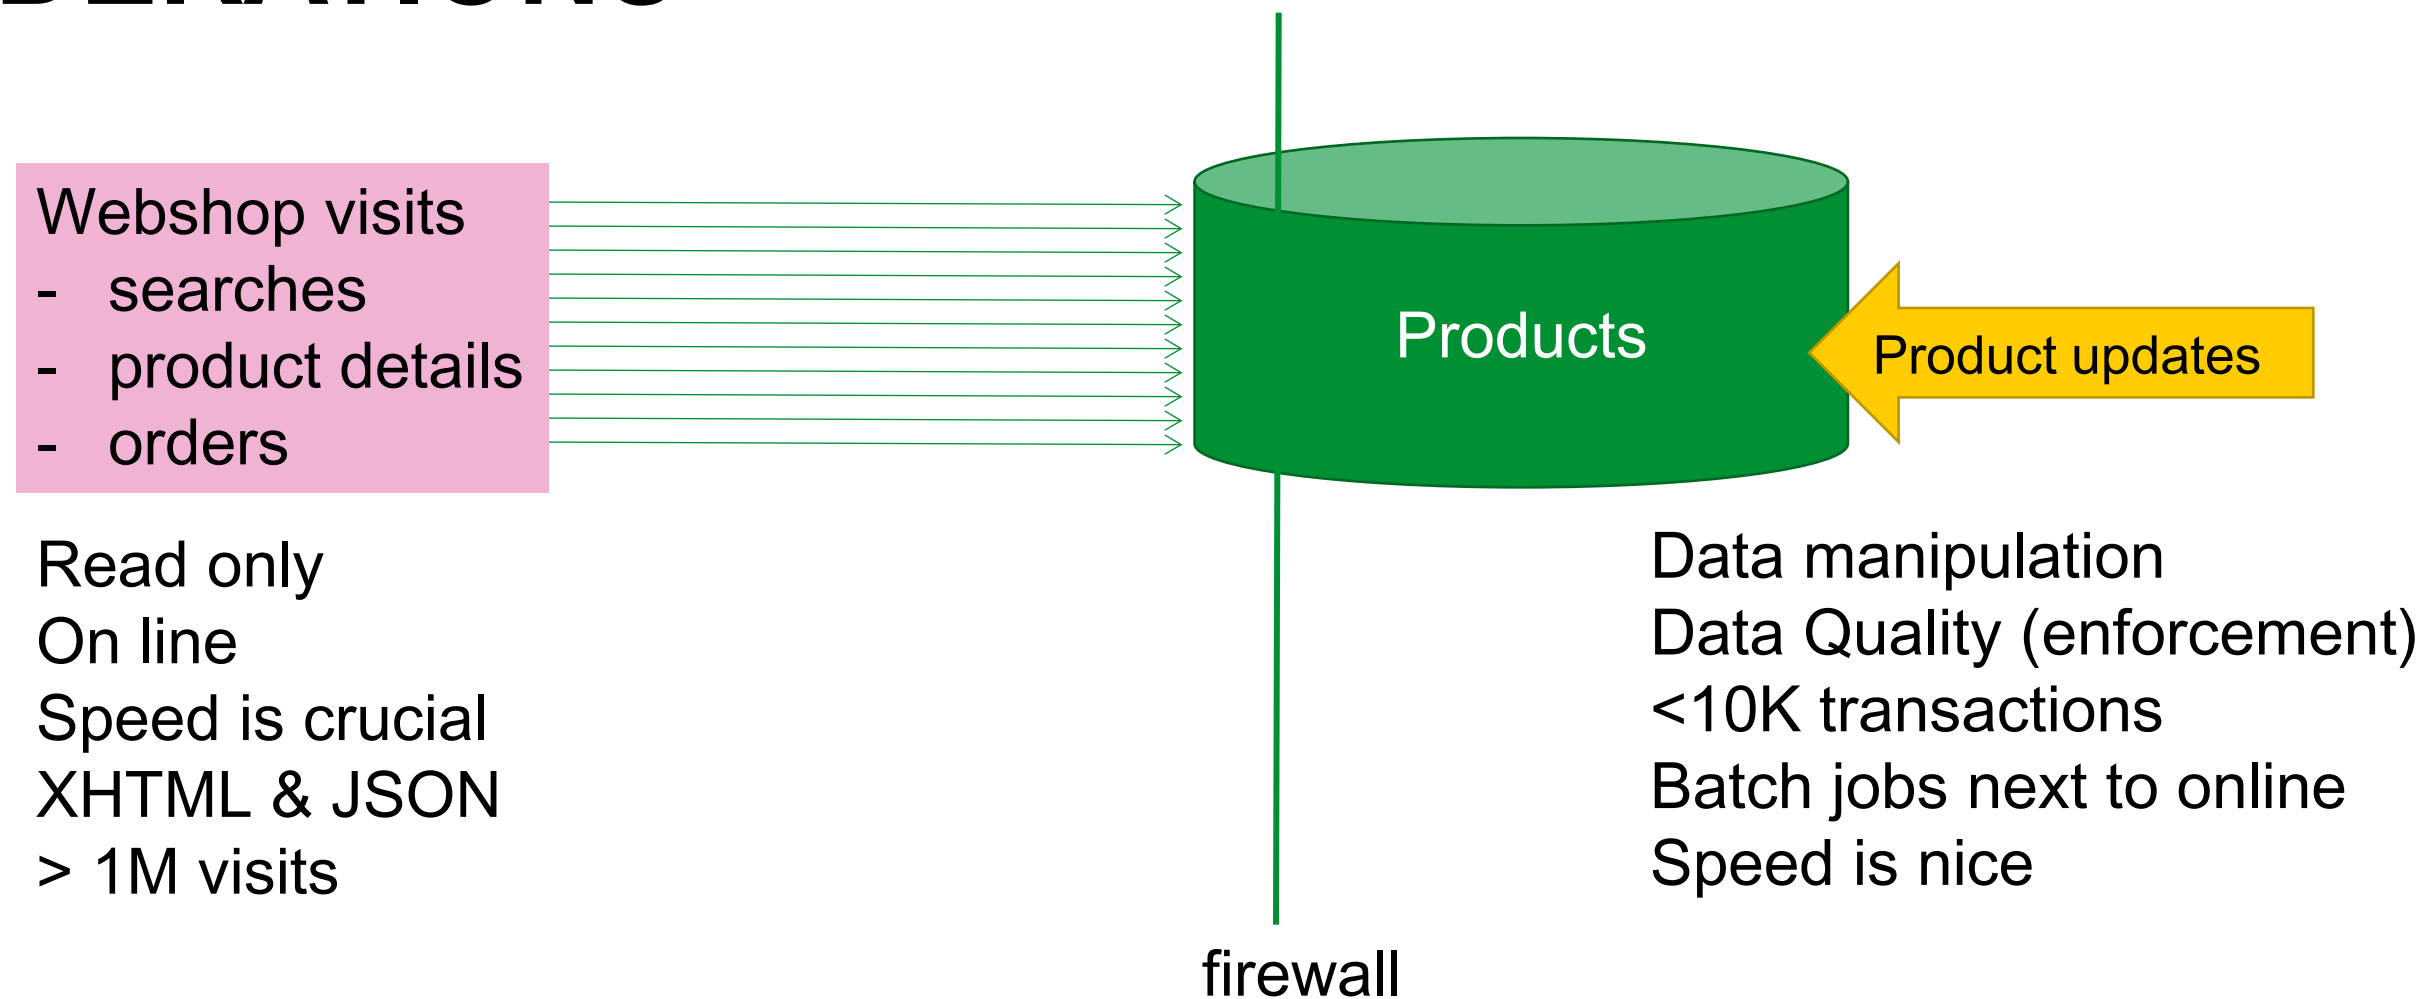
- Not all use cases require ACID (or can afford it)
 - Read only (product catalog for web shops)
 - Inserts only and no (inter-record) constraints
 - Big Data collected and “dumped” in Data Lake (Hadoop) for subsequent processing
 - High performance demands
- Not all data needs structured formats or structured querying and JOINS
 - Entire documents are stored and retrieved based on a single key
- Sometimes – scalable availability is more important than Consistency – and ACID is sacrificed
 - CAP-theorem states: Consistency [across nodes], Availability and Partition tolerance can not all three be satisfied

USE CASE: WEBSHOP

- Webshop – 1M visitors per day
- Product catalog consists of millions of records
 - The web shop presents: product description, images, reviews, pricing details, related offerings, stock status
- Products are added and updated and removed every day
 - Although most products do not change very frequently
 - Some vendors do bulk manipulation of product details

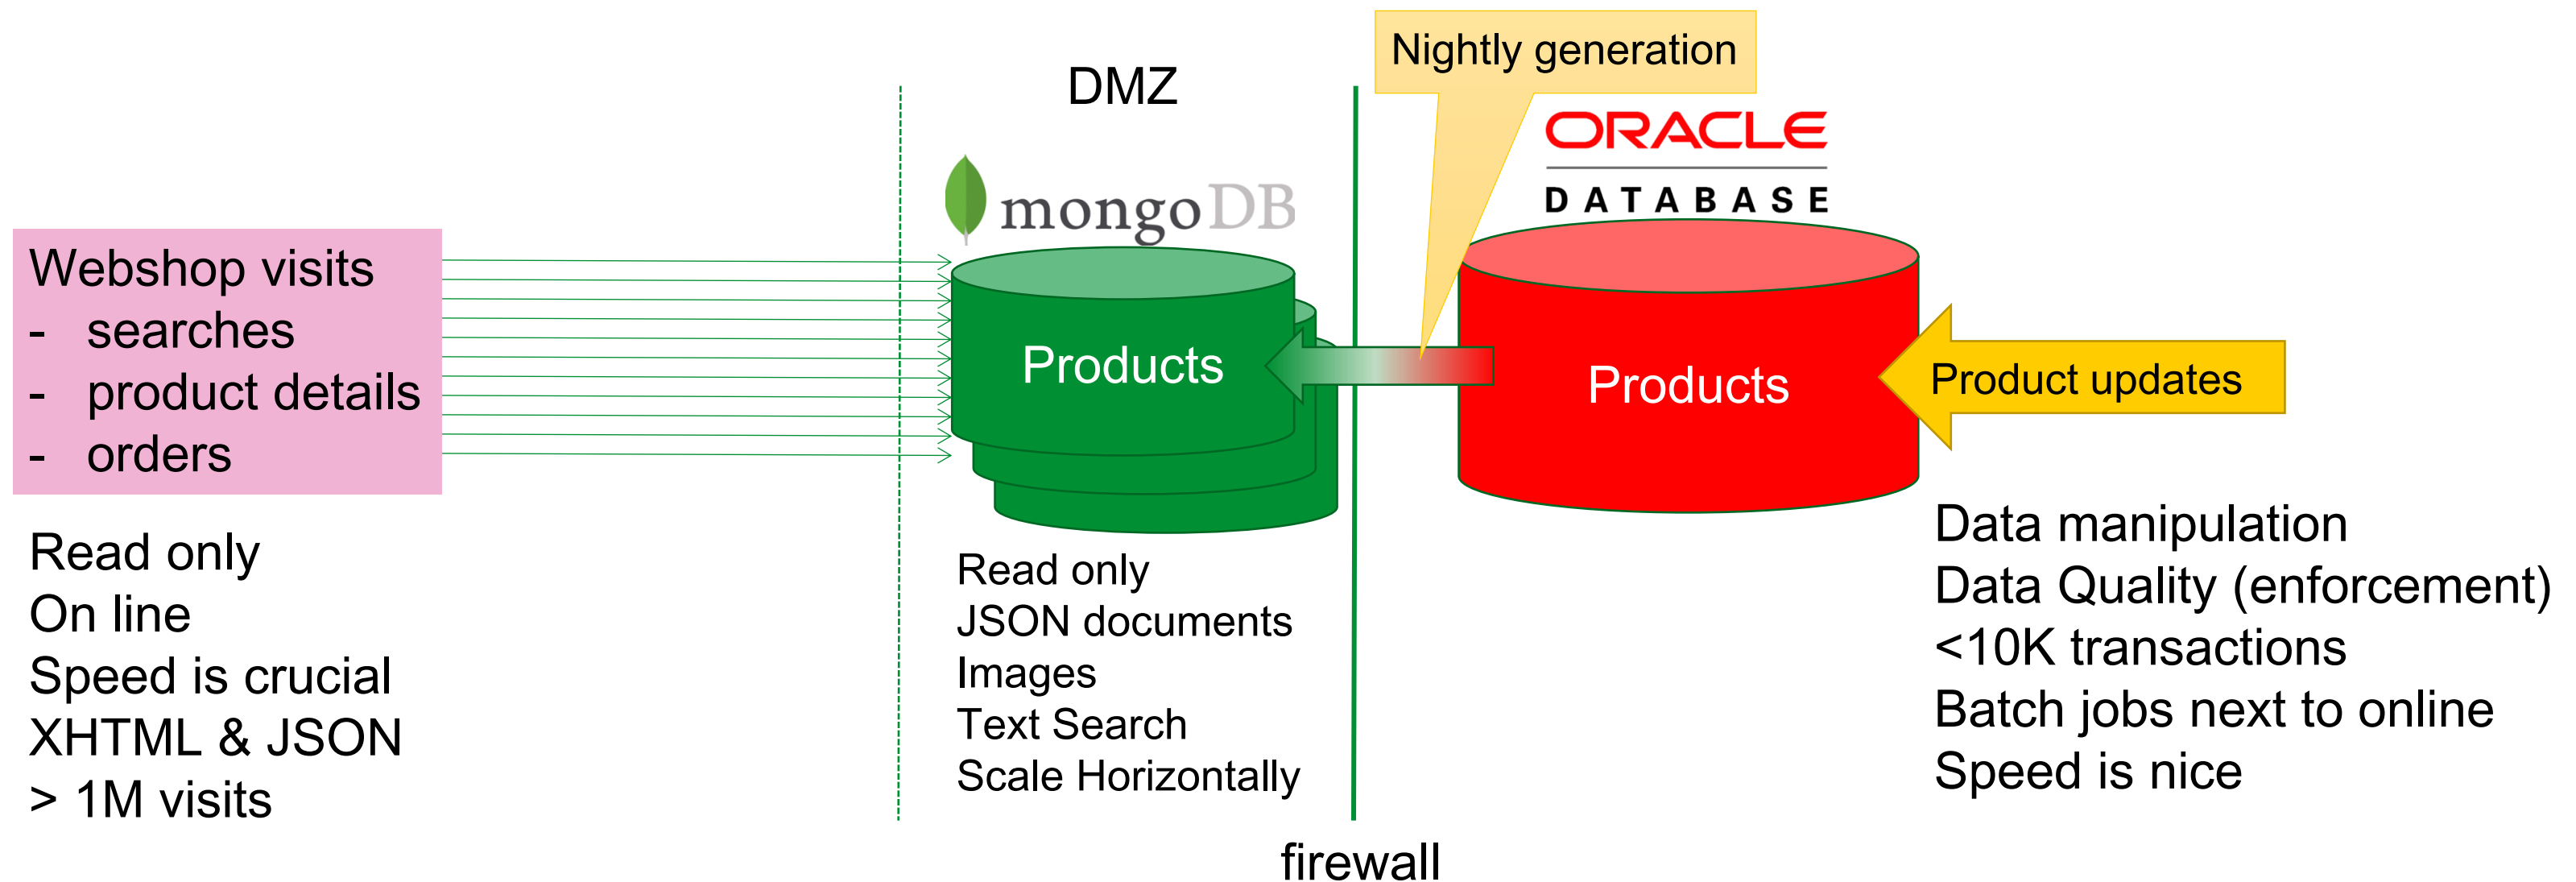


USE CASE: WEBSHOP TECHNOLOGY AND ARCHITECTURE CONSIDERATIONS



USE CASE: WEBSHOP

TECHNOLOGY AND ARCHITECTURE CONSIDERATIONS



A FEW OBSERVATIONS

- Redundancy
 - for performance, scalability, availability
- Commodity hardware for scale out architecture
- Data (documents) prepared for specific application usage (super MV)
 - Format
 - Search paths
- Time lag in data modifications is accepted
 - Perhaps even desirable to offer a consistent view to webshop users
- Transactions in one engine
 - ACID still in place
- Transactions do not compete with reads
 - vying for the same physical resources

NOSQL AND BASE

- NoSQL arose because of performance and scalability challenges with Web Scale operations
- NoSQL is a label for a wide variety of databases that aspect of a true relational database
 - ACID-ness, SQL, relational model, constraints
- The label has been used since 2009
 - Perhaps NoREL would be more appropriate
- Some well known NoSQL products are
 - Cassandra, MongoDB, Redis, CouchDB, ...
- BASE as alternative to ACID:
 - **b**asically **a**vailable, **s**oft state, **e**ventually **c**onsistent
(after a short duration)



TYPICAL FOR NOSQL

- Focus on speed, availability and scalability
 - Horizontal scale out – distributed with load balancing and fail-over
- No (predefined) Data Structure
- Integrity primarily protected by application logic
- Open Source (most offerings are, not all: MarkLogic)
- Close(r) attention for how the data is used
 - Application oriented data format and search paths and specialized database per application (microservice, capability)
 - Similar to the switch from SOA to API/Microservice
- Reads (far) more relevant than writes
- Data redundancy & denormalization
- No data access through SQL

NOSQL MEANS: NO DATA ACCESS THROUGH SQL

- However
 - Data Professionals and Developers speak SQL
 - Reporting, Dashboarding, ETL, BI tools speak SQL
- There is no common query language across NoSQL products

The screenshot shows the CDATA website with a navigation bar at the top containing links for PRODUCTS, DOWNLOAD, SUPPORT, ORDER, and COMPANY. On the right, there's a link to "SEE THE WORLD AS A DATABASE" and contact information for Chat, Cart, and a phone number (800.235.7250).

JDBC Drivers

Easy-to-use JDBC Drivers with powerful Enterprise-level features

Straightforward access to live Application, Database, and WebAPI data through standard Java database connectivity.

- Unparalleled performance and scalability.
- Simple JDBC / SQL access to live data with full read/write access.
- Fully-integrated access to live data from popular BI, Analytics, and Reporting Tools.
- Integration into popular IDEs like Eclipse, IntelliJ, and NetBeans.

On the right side of the page, there's a graphic showing various BI and Analytics tools (COGNOS, Tableau, KNIME, Cold Fusion, Pentaho, ETL, BI Tools, Analytics, Custom Apps) connected to a central database icon, with a Java logo at the bottom right.

JDBC Drivers for:

NoSQL & Big Data

Amazon DynamoDB	Amazon SimpleDB	Google BigQuery
Azure Storage	Cassandra	1010data
Elasticsearch	HPCC Systems	Database.com
MongoDB	Couchbase	HBase

NO DATA ACCESS THROUGH SQL

- However
 - Data Professionals and Developers speak SQL
 - Reporting, Dashboarding, ETL, BI tools speak SQL
- There is no common query language across NoSQL products
- Attempts from many vendors to create drivers that translate SQL statements into NoSQL commands for the specific target database
 - To protect existing investments in SQL – skills, tools, applications, reports, ..

Oracle Big Data SQL
One Fast, Secure Query over All your Data



cdata SEE THE WORLD AS A DATABASE
PRODUCTS ▾ DOWNLOAD ▾ SUPPORT ▾ ORDER ▾ COMPANY ▾ Chat Cart 800.235.7250

JDBC Drivers

Easy-to-use JDBC Drivers with powerful Enterprise-level features

Straightforward access to live Application, Database, and WebAPI data through standard Java database connectivity.

- Unparalleled performance and scalability.
- Simple JDBC / SQL access to live data with full read/write access.
- Fully-integrated access to live data from popular BI, Analytics, and Reporting Tools.
- Integration into popular IDEs like Eclipse, IntelliJ, and NetBeans.

JDBC Drivers for:

NoSQL & Big Data


Amazon DynamoDB	Amazon SimpleDB	Google BigQuery
Azure Storage	Cassandra	1010data
Elasticsearch	HPCC Systems	Database.com
MongoDB	Couchbase	HBase

TYPES OF NOSQL DATABASES

- Key-Value (from in-memory cache to hardware appliance – Memcached, Hazelcast, Coherence, Redis, Oracle NoSQL/Berkeley DB)
- Wide Column Store (Cassandra, Google BigTable, HBase)
- Document (JSON, XML, YAML and binary forms – MongoDB, CouchDB, Couchbase, MarkLogic)
- Graph (networks of data, triplets, RDF stores – Neo4J)
- Miscellaneous: object, tabular, tuple, elastic search index, Hadoop (hdfs) style, Kafka Topic

(LEADING) NOSQL DATABASE PRODUCTS

- MongoDB is (one of) the most popular (by any measure)
- Cloud (only):
 - Google BigTable,
 - AWS Dynamo
- Cache (in memory)
 - ZooKeeper, Redis, Memcached, ...
- Hadoop/HDFS
- Oracle NoSQL (fka Berkeley DB)





HISTORY OF MONGODB

- 10gen – startup from 2007 in New York City
 - Developed database as component in Platform as a Service product
- Initial release of MongoDB: 2009
 - 10gen changed its name to MongoDB Inc. in 2013
 - Offers enterprise support, MongoDB Atlas cloud service, MongoDB Compass, Connectors
- Current stable release: 3.4.2 (3.5.2 is in preview)
- Open Source - GNU Affero General Public License and the Apache License.
- MongoDB == *Humongous Database*
- Evolution of MongoDB is still pretty fast
 - Some crucial enterprise database capabilities were added fairly recently or are still lacking
- MongoDB is number one NoSQL database in terms of popularity




CORE ARCHITECTURE

- MongoDB stores JSON documents in a binary format (BSON)
 - Documents are stored in collections (similar to row or records in tables)
- Interaction is through JavaScript
 - taking the place of SQL for DML and query
- Written in C++
- Has the V8 JavaScript engine included
- Runs on Mac OS X, Windows, Solaris, and most flavors of Linux
- Source code is on GitHub
- Has GeoSpatial and Tekst indexes and search capabilities
- MongoDB, Inc. officially supports drivers for C, C++, C#, Erlang, Java, Node.js, JavaScript, Perl, PHP, Python, Scala, and Ruby




ARCHITECTUUR, INSTALLATIE, CONFIGURATIE, OPSLAG, SHELL




- Mongo draait op windows en linux.
- Windows install met msi file
- linux install met tgz file.
- Users en directories vrij te kiezen
- Mongod --dbpath db1-- port 27001
- Default path=\data\db, default port=27017
- Mongo [-- port hostname:27001] [--shell some_json_file]
- Dbs, collections

IMPORT EXPORT




- Mongoimport
- Json
- Csv, tsv, maar hebben een descriptor file nodig
- mongoexport -h -d -c -o -q
- Leesbare output, minder geschikt voor backups

CRUD




- show dbs
- use db
- db.<collection>.find()
- db.<collection>.find().pretty()
- select: db.find({key:value,key:value})
- present: db.<collection>.find({key:value},{key:1,key:1,_id:0})
- explain: db.find({key:value,key:value}).explain("executionStats")
- db.companies.insert({Company:"AMIS",Event:"MongoDB Workshop"})
- db.companies.update({Company:"AMIS"},{\$set: {Location:"Nieuwegein"}})
- Geen commit! Ofwel, autocommit.

CRUD




- `db.companies.remove({"Company" : "AMIS"})`
- `$unset` -> Om een veld uit een document te verwijderen.
- `$inc` -> om een numeriek veld te verhogen met een bepaalde waarde. Als het veld niet bestaat wordt het toegevoegd (met de waarde van de verhoging)
- `$push` -> Om een waarde aan een array toe te voegen.

INDEXING, EXPLAIN




- `_id` altijd geïndexeerd en unique
- Indexen helpen bij zoeken
- Indexen mogen compound zijn, dwz op meerdere keys gecombineerd, bv index op woonplaats+leeftijd
- `db.user.find().sort({ woonplaats:1, leeftijd: -1 })` De -1 betekent aflopend geïndexeerd.
- Resultaten rechtstreeks uit index indien mogelijk (fat indexes in SQL)
- `db.user.find({woonplaats:"Zoetermeer"}).explain("executionStats")` geeft een goed overzicht van het gevolgde plan

REPLICA SETS




- Replica sets zijn gedupliceerde collections
- Beveiliging tegen uitval
- De members besluiten onderling welk member primary is. Alleen daarop kan DML plaatsvinden
- Altijd oneven aantal members, eventueel mbv een arbiter
- Secondary members kunnen wel voor select gebruikt worden, als je eerst slaveOk() ingeeft
- Mogen op dezelfde host draaien, mogen ook op verschillende hosts draaien.
- writeConcern: Je kan expliciet aangeven op hoeveel members je DML uitgevoerd moet zijn voor je je prompt terug krijgt.

SHARDING




- Shard = splinter
- Range based partitioning
- Altijd index nodig op de shard key. Hoeft niet unique te zijn.
- Iedere shard is een proces. Bij voorkeur ieder proces zijn eigen host.
- Scale out proces
- Metadata staat in aparte config db. Typically 3 stuks.
- Clients verbinden met een mongos server. Dat is een instance zonder database die uit de config haalt waar de data zich bevindt.
- Iedere shard kan bestaan uit replica sets.

SHARDING




- Data verdeelt zichzelf over chunks.
- Chunks kunnen splitten wanneer ze te groot worden (typically 100MB maar is aan te passen naar wens)
- Chunks kunnen migraten naar een andere shard. Mongo streeft ernaar elke shard evenveel chunks te laten hebben.
- Tijdens migration kan DML gewoon doorgaan.

BACKUP & RESTORE

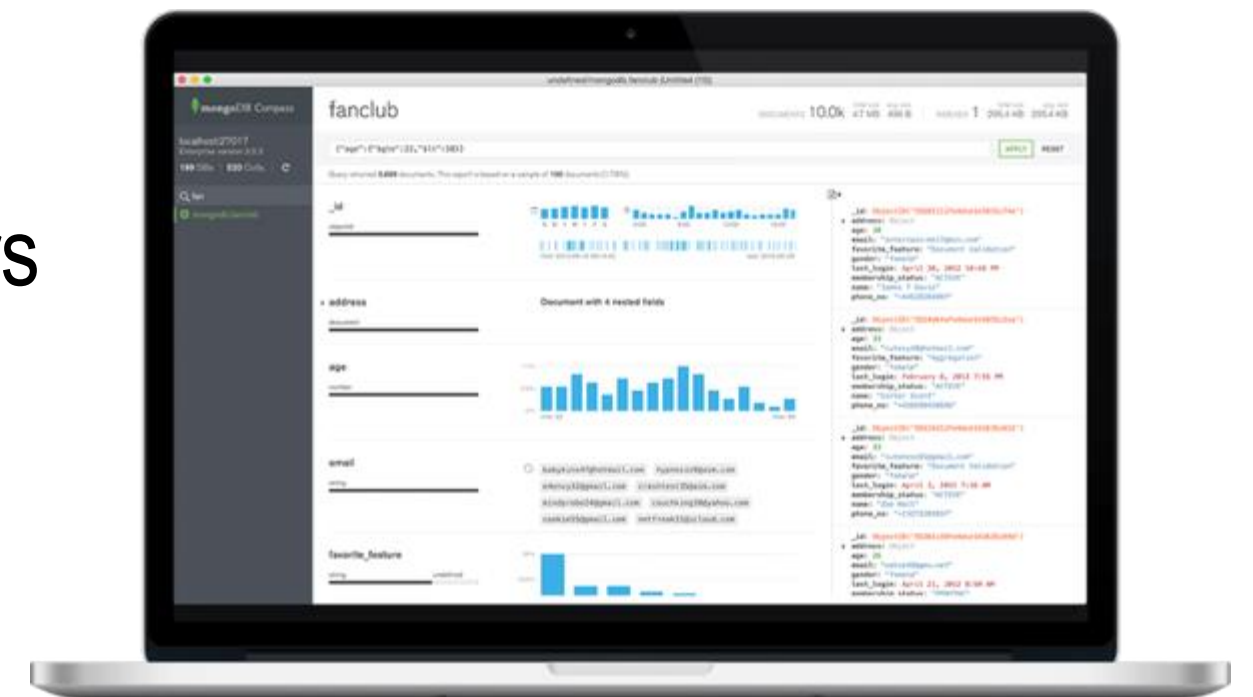


- Backup met mongodump
- Creëert een dump directory met daaronder een directory per database en plaatst daarin 1 of meerdere .bson files, een binair formaat.
- Kan vanaf ieder member van een replica set. Bij voorkeur eentje die het niet druk heeft.
- Mongo kent geen transactions, dus backup is niet consistent. Het is maar afwachten wat erin staat als er veel DML plaatsvindt.
- De tegenhanger heet mongorestore.
- Restore van bestaande database lukt niet: duplicate key op _id

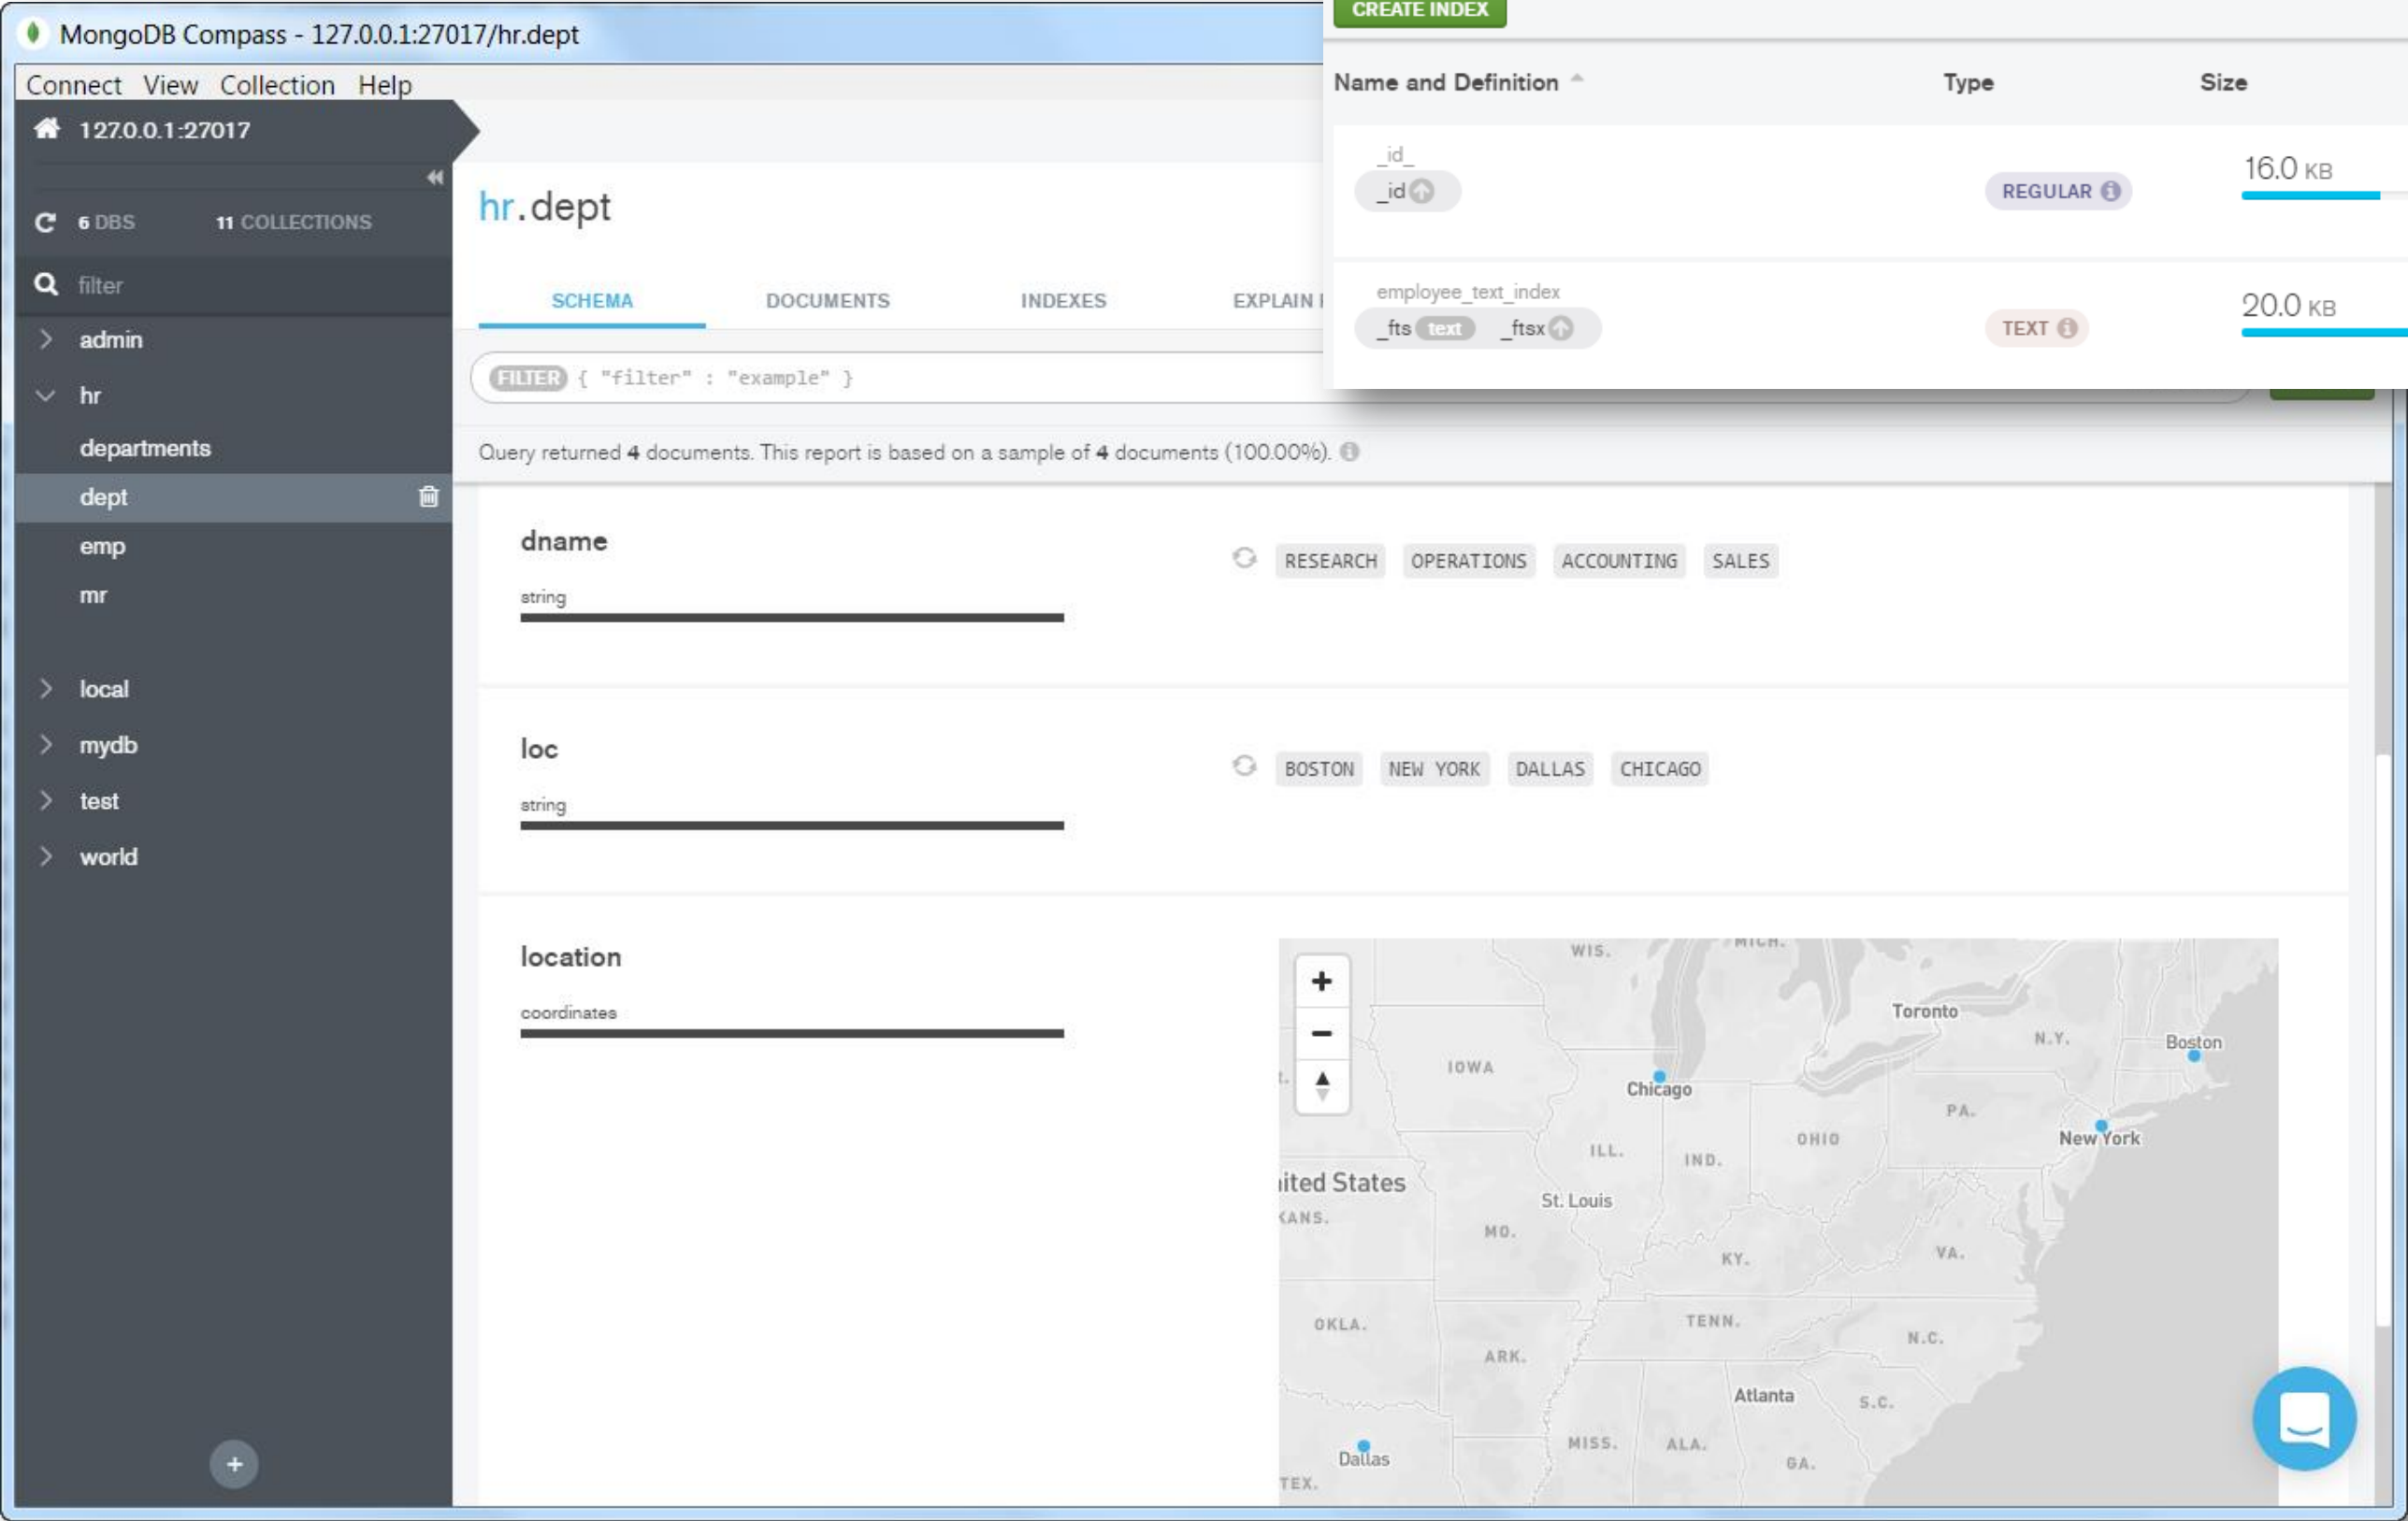
MONGODB COMPASS (AKA NOSQL DEVELOPER)



- The GUI for MongoDB
 - Visually explore data
 - Run ad hoc queries in seconds.
 - Interact with your data with full CRUD functionality
 - View and optimize your query performance.
- Separately installed product
 - On Windows, OS X, Linux, ...
- Can run against local and remote MongoDB servers



MONGODB COMPASS



hr.emp

DOCUMENTS 14 TOTAL SIZE 6.6KB AVG. SIZE 485B INDEXES 2 TOTAL SIZE 36.0KB AVG. SIZE 18.0KB


SCHEMA DOCUMENTS INDEXES EXPLAIN PLAN VALIDATION

CREATE INDEX

Name and Definition ^	Type	Size	Usage	Properties	Drop
<u>_id_</u> _id_ ↑	REGULAR ⓘ	16.0 KB	0 since Wed Mar 08 2017	UNIQUE ⓘ	
employee_text_index _fts text _ftsx ↑	TEXT ⓘ	20.0 KB	0 since Wed Mar 08 2017	COMPOUND ⓘ	🗑

PROGRAMMING AGAINST MONGODB

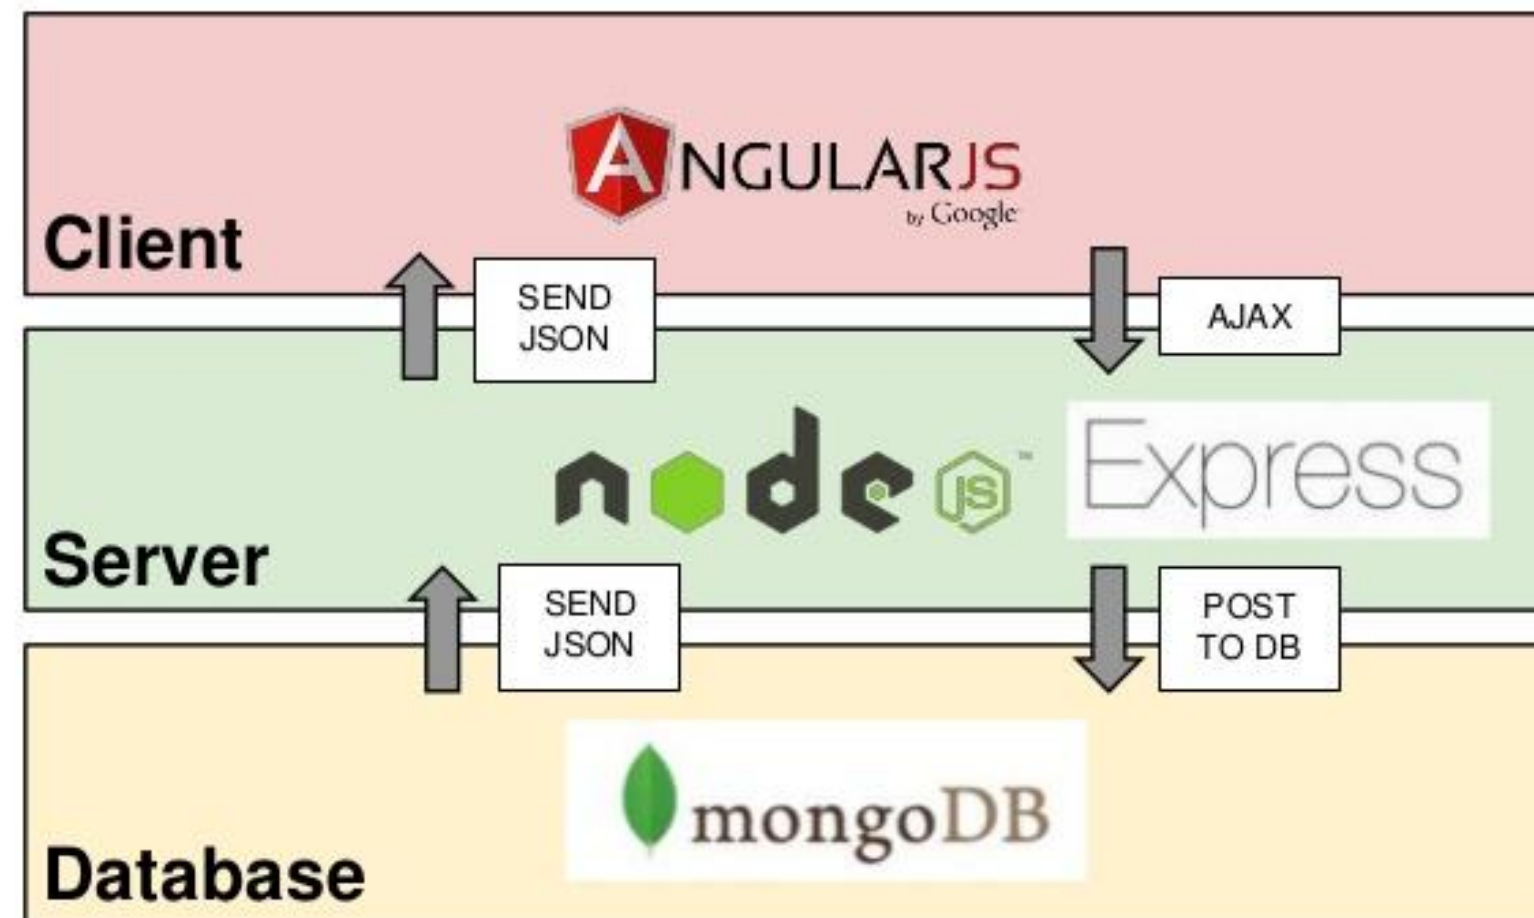
- MongoDB officially supports drivers for C, C++, C#, Erlang, *Java*, *Node.js*, JavaScript, Perl, PHP, Python, Scala, and Ruby



- Note: there is also a Connector for Hadoop
- Several REST APIs for MongoDB are available from the community
 - RESTHeart (Java), Eve (Python), Crest (Node), AMID, Kule, DreamFactory


MEAN STACK

- Modelled after LAMP
- End-to-End JavaScript/JSON
- Term coined in 2013 (by MongoDB)




INTERACTING WITH MONGODB FROM NODE.JS

- MongoDB Node.js Driver 2.0
 - Support for ECMAScript 6.0 – Promises for asynch
 - <http://mongodb.github.io/node-mongodb-native/2.0/>
- Using mongodb in a NodeJS application
 - `npm install mongodb --save`




```
package.json x
1 {
2   "name": "mongodb-nodejs-client",
3   "version": "1.0.0",
4   "author": "Lucas J. J. J.",
5   "license": "ISC",
6   "dependencies": {
7     "mongodb": "^2.2.24"
8   }
9 }
```

- `var MongoClient = require('mongodb').MongoClient;`



CONNECTING TO SERVER



```
var MongoClient = require('mongodb').MongoClient;

var mongodbHost = '127.0.0.1';
var mongodbPort = '27017';
var mongodbDatabase = 'world';


var url = 'mongodb://' + mongodbHost + ':' + mongodbPort + '/' + mongodbDatabase;

MongoClient.connect(url, function(err, db) {
  console.log("Connected correctly to server.");

  // DO YOUR THING WITH MONGODB

  db.close();
  console.log("Connection to database is closed.");
}) //connect()
```

RETRIEVE DATA (1)



```
var MongoClient = require('mongodb').MongoClient;
var mongodbHost = '127.0.0.1';
var mongodbPort = '27017';
var mongodbDatabase = 'world';
var url = 'mongodb://' + mongodbHost + ':' + mongodbPort + '/' + mongodbDatabase;


MongoClient.connect(url, function(err, db) {
  console.log("Connected correctly to server.");
  db.collection('countries').find({}, {"sort": [{"area",-1}]})
    .limit(20).toArray(function(err, results){
      console.log("Name of Country Four " + results[3].name+ " and size: " + results[3].area);

      // using cursors to retrieve data sets in a controlled fashion
      // note: cursor implements NodeJS Stream - results can be piped
      var cursor = db.collection('countries').find({"continent":"Asia"}, {"sort": "name"});

      cursor.count(function(err, count){
        console.log("Country count in Asia: "+count);
      });

      cursor.each(function(err, country){
        console.log(country.name);
      })//each cursor
```



RETRIEVE DATA (2)



```
var aggquery = [  {$sort: {area : -1}}
                  ,  {$group:{ _id: '$continent'
                                , largestCountry : {$first: "$name"}
                              }}
                ];
var aggcursor = db.collection('countries').aggregate(aggquery);
aggcursor.each(function(err, result){
  if (err) {
    console.log(err);
  } else if (result)
    console.log(JSON.stringify(result));
}) //each aggcursor

var ccursor = db.collection('countries').find({});
// the cursor returned from find and aggregate implements a NodeJS (readable) Stream
ccursor.on('data', function(doc) {
  console.log(doc);
});
ccursor.once('end', function() {
  console.log("Out of countries. Time to move on");
});
```

CONNECTING TO SERVER – ECMA 6



```
var MongoClient = require('mongodb').MongoClient;
var assert = require('assert'),    co = require('co');


var mongodbHost = '127.0.0.1';
var mongodbPort = '27017';
var mongodbDatabase = 'world';
var url = 'mongodb://' + mongodbHost + ':' + mongodbPort + '/' + mongodbDatabase;

co(function*() {
  // Use connect method to connect to the Server
  var db = yield MongoClient.connect(url);

  console.log("Connected correctly to server");
  // DO YOUR THING WITH MONGODB


  // Close the connection
  db.close();
}).catch(function(err) {
  console.log(err.stack);
});
```

RETRIEVE DOCUMENTS – ECMA 6




```
co(function*() {  
  // Use connect method to connect to the Server  
  var db = yield MongoClient.connect(url);  
  
  console.log("Connected correctly to server");  
  // find top 20 countries by size  
  var results = yield db.collection('countries').find({},  
    {"sort": [{"area",-1}]).limit(20).toArray();  
  console.log("Country One " +JSON.stringify(results[0]));  
  console.log("Name of Country Four " +results[3].name  
    + " and size: " +results[3].area);  
  
  // use cursor to get the country count  
  var cursor = db.collection('countries').find({"continent":"Asia"},  
    {"sort": "name"});  
  
  var count = yield cursor.count();  
  console.log("Country count in Asia: "+count);  
  
  while (yield cursor.hasNext()){  
    var country = yield cursor.next();  
    console.log(country.name);  
  }  
}
```

RETRIEVE DOCUMENTS (2) – ECMA 6




```
// the largest country per continent
var aggquery = [  {$sort: {area : -1}}
                  ,  {$group:{ _id: '$continent'
                               , largestCountry : {$first: "$name"}
                              }}
                ];
var aggcursor = db.collection('countries').aggregate(aggquery);
while (yield aggcursor.hasNext()){
    var result = yield aggcursor.next();
    console.log(JSON.stringify(result));
}
// Close the connection
db.close();
}).catch(function(err) {
    console.log(err.stack);
}); //co
```


CREATE DOCS IN MONGODB – ES6




```
// define two documents, any structure you like
var doc = {
  "continent" : "Europe",
  "children" : [ {"name":"Belgium"}, {"name":"Luxemburg"} ],
  "someVariable" : 19123, "andmore" : "2716NK" };
var doc2 = {
  "continent" : "Asia", "name" : "Johnny",
  "nrs" : [ {"name":"China"}, {"name":"India"}, {"name":"Buthan"} ],
  "tree": {"branch": {"twig":{"leaf": 1}}}
};
var nameOfCollection = "myDocs"
var result = yield db.collection(nameOfCollection).insertMany([doc,doc2]);
console.log(">> "+result.insertedCount
           +" documents created into collection "+nameOfCollection);
...
```

CREATE & UPDATE DOCS – ES6



```
// define two documents, any structure you like
var doc = {...}, doc2 = {...} , nameOfCollection = "myDocs";
var result = yield db.collection(nameOfCollection).insertMany([doc,doc2]);
console.log(">> "+result.insertedCount +" documents created");
var cursor = db.collection(nameOfCollection).find();
while (yield cursor.hasNext()){
    var doc  = yield cursor.next();
    console.log("Document: " +JSON.stringify(doc));
}// while cursor
result = yield db.collection(nameOfCollection).updateOne(
    {"tree.branch.twig.leaf":1}, {$set: {name: "Hank", city:"Nieuwegein",
    "location.province":"Utrecht",
    "location.country":"The Netherlands",
    "tree.stem":5}});
console.log(">> Updated "+result.modifiedCount+" document(s)");
cursor = db.collection(nameOfCollection).find();
while (yield cursor.hasNext()){
    var doc  = yield cursor.next();
    console.log("Document: " +JSON.stringify(doc));
}// while cursor
```

CREATE, UPDATE AND DELETE – ES6



```
// define two documents, any structure you like
var doc = {...}, doc2 = {...} , nameOfCollection = "myDocs";
var result = yield db.collection(nameOfCollection).insertMany([doc,doc2]);
console.log(">> "+result.insertedCount +" documents created");
var cursor = db.collection(nameOfCollection).find();
while (yield cursor.hasNext()){
    var doc  = yield cursor.next();
    console.log("Document: " +JSON.stringify(doc));
} // while cursor
result = yield db.collection(nameOfCollection).updateOne(...);
...
result = yield db.collection(nameOfCollection).deleteMany({});
console.log(">> Deleted "+ result.deletedCount+" documents ");
// execute command to drop the collection
yield db.command({drop:nameOfCollection});
console.log(">> Dropped collection "+nameOfCollection);
// Close the connection
db.close();
console.log("Connection to database is closed.");
```


INTERACTING WITH MONGODB FROM JAVA

- MongoDB Java Driver 3.x
 - Supports Synchronous and Asynchronous operations
 - <https://docs.mongodb.com/ecosystem/drivers/java/>
- Using MongoDB Driver in a Java Application
 - Add dependency in Maven pom.xml file


```
<dependencies>
  <dependency>
    <groupId>org.mongodb</groupId>
    <artifactId>mongo-java-driver</artifactId>
    <version>3.4.0</version>
  </dependency>
</dependencies>
```

- Library JAR file mongo-java-driver-3.4.0.jar

Name	Type	Size
junit-3.8.1	Executable Jar File	119 KB
mongo-java-driver-3.4.0	Executable Jar File	1.634 KB



CONNECTING TO SERVER




```
package nl.amis.mongodb.countries;

import java.net.UnknownHostException;
import com.mongodb.MongoClient;
import com.mongodb.client.MongoDatabase;

public class AppConnect {
    public static void main(String[] args) throws UnknownHostException {
        String mongodbHost = "127.0.0.1";
        Integer mongodbPort = 27017;
        String mongodbDatabase = "world";

        MongoClient mongoClient = new MongoClient(mongodbHost, mongodbPort);
        try {
            System.out.println("Connected to server; now hook into database " + mongodbDatabase);
            MongoDatabase db = mongoClient.getDatabase(mongodbDatabase);
            System.out.println("Connected to database " + mongodbDatabase + " successfully");
        } finally {
            mongoClient.close();
            System.out.println("Closed mongodb connection");
        }
    }
}
```


RETRIEVE DATA



```
...
import org.bson.Document;
import com.mongodb.client.MongoCollection;


public class AppConnect {
    public static void main(String[] args) throws UnknownHostException {
        ... Connect to world database

// get all documents from the countries collection
MongoCollection<Document> countries = db.getCollection("countries");
System.out.println("Got collection countries successfully");

System.out.println("Number of countries: " + countries.count());

System.out.println("All countries - name and continent");
for (Document doc : countries.find()) {
    // System.out.println(doc.toJson());
    System.out.println(doc.get("name") + " (" + doc.get("continent") + ")");
} //for
```

RETRIEVE DATA (2)




```
import com.mongodb.client.model.Sorts;
import com.mongodb.client.model.Filters;
import com.mongodb.Block;

public class AppConnect {
    public static void main(String[] args) throws UnknownHostException {
        ... Connect to world database

        System.out.println( "Retrieve all countries in Asia order alphabetically
                               by name and display full country document");

        countries.find(Filters.eq("continent", "Asia"))
            .sort(Sorts.ascending("name"))
            .forEach(new Block<Document>() {
                @Override
                public void apply(final Document document) {
                    System.out.println(document.toJson());
                }
            });
    }
}
```


RETRIEVE DATA - AGGREGATION



```
...
Block<Document> printBlock = new Block<Document>() {
    @Override
    public void apply(final Document document) {
        System.out.println("Country: " + document.toJson());
    }
};

...
System.out.println("List largest country in each continent");
countries.aggregate(
    Arrays.asList(
        Aggregates.sort(Sorts.descending("area")),
        Aggregates.group("$continent", Accumulators.first("Largest Country", "$name"))
    )
    .forEach(printBlock);
```

RETRIEVE DATA – AGGREGATION (2)




```
System.out.println("List name in uppercase,sorted by size, delta between birthdate and deathrate, "
                    +" population density for countries in Oceania");

BsonArray subArgs=new BsonArray();
    subArgs.add(new BsonString("$birthrate"));
    subArgs.add(new BsonString("$deathrate"));
BsonArray divArgs=new BsonArray();
    divArgs.add(new BsonString("$population"));
    divArgs.add(new BsonString("$area"));

countries.aggregate
(Arrays.asList(
    Aggregates.match(Filters.eq("continent", "Oceania")),
    Aggregates.sort(Sorts.descending("area")),
    Aggregates.project( Projections.fields(
        Projections.excludeId(),
        Projections.include("continent"),
        Projections.computed("name",Projections.computed("$toUpper",  "$name")),
        Projections.computed("populationGrowthRate",Projections.computed("$subtract",  subArgs)),
        Projections.computed("populationDensity",
            Projections.computed("$trunc", Projections.computed("$divide",  divArgs)))
    ))
))
.forEach(printBlock);
```


JAVA: CREATE DOCUMENTS IN MONGODB



```
....
String nameOfCollection = "myDocs";
MongoCollection<Document> coll = db.getCollection(nameOfCollection);
Document doc1 = new Document("continent", "Europe")
    .append("nrs", Arrays.asList(new Document("name", "Belgium"),
    new Document("name", "Luxemburg")))
    .append("someVariable", 19123)
    .append("andmore", "kl;jdsfakhfsdahjfsdasjbsdahjbsdahgvsahjkZl;po");
Document doc2 = new Document("continent", "Asia").append("name", "Johnny")
    .append("nrs", Arrays.asList(new Document("name", "China"),
    new Document("name", "India"), new Document("name", "Buthan")))
    .append("tree", new Document("branch",
    new Document("twig", new Document("leaf", 1))));
List<Document> documents = new ArrayList<Document>();
documents.add(doc1);
documents.add(doc2);

coll.insertMany(documents);
System.out.println(">> Documents created into collection");
coll.find().forEach(printBlock);
```



JAVA: UPDATE AND DELETE DATA



```
...
coll.updateOne(
    Filters.eq("tree.branch.twig.leaf", 1),
    Updates.combine(
        Updates.set("name", "Hank"), Updates.set("city", "Nieuwegein"),
        Updates.set("location.province", "Utrecht"),
        Updates.set("location.country", "The Netherlands"),
        Updates.set("tree.stemp", 5),
        Updates.currentDate("lastModified")
    )
);
System.out.println(">> Updated one document in collection " + nameOfCollection);
coll.find().forEach(printBlock);

DeleteResult dr = coll.deleteMany(Filters.exists("_id"));
System.out.println(">> Deleted " + dr.getDeletedCount() + " documents from collection "
    + nameOfCollection);
db.runCommand(new Document("drop", nameOfCollection));
System.out.println(">> Dropped collection " + nameOfCollection);
```

INTERACTING WITH MONGODB FROM JAVA – JSON ↔ OO MAPPING



- Libraries for mapping between BSON/JSON documents and Java Objects
 - Morphia - <http://mongodb.github.io/morphia/1.3/getting-started/quick-tour/>
- 💡 • Jongo – Query in Java as in the MongoDB shell - <http://jongo.org/>


```
Jongo jongo = new Jongo(db);
MongoCollection friends = jongo.getCollection("friends");

MongoCursor<Friend> all = friends.find("{name: 'Joe'}").as(Friend.class);
Friend one = friends.findOne("{name: 'Joe'}").as(Friend.class);
```

HANDS ON WORKSHOP

- Install MongoDB
- Create Database, Create Collections and Documents
- Retrieve Data
- Import Database
- More configuration and administration
- Development: interacting with MongoDB from
 - Compass
 - Java
 - Node.js





- **Resources:** <https://github.com/lucasjellema/sig-nosql-mongodb>

- **Blog:** technology.amis.nl

- **Email:** pom.bleeksma@amis.nl , lucas.jellema@amis.nl

-  : [@lucasjellema](https://twitter.com/lucasjellema)

-  : [lucas-jellema](https://www.linkedin.com/in/lucas-jellema)

- : www.amis.nl, info@amis.nl



+31 306016000
Edisonbaan 15,
Nieuwegein